

# Week 00a Basic Handout: ML Foundations - Learning from Data

Machine Learning for Smarter Innovation

## 1 Week 00a Basic Handout: ML Foundations - Learning from Data

### 1.1 For Students With: No prior ML knowledge, basic programming concepts

#### 1.2 Overview

This handout introduces machine learning fundamentals without requiring mathematical background. Focus on concepts, real-world examples, and practical understanding.

---

### 1.3 Part 1: The Big Idea - Machines That Learn

#### 1.3.1 What is Machine Learning?

Machine learning means teaching computers to find patterns in data instead of programming explicit rules.

**Traditional Programming:** - Programmer writes: IF spam word THEN mark as spam - Works for simple cases - Breaks with edge cases

**Machine Learning:** - Computer discovers: "These 500 emails are spam, these 500 are not spam" - Learns patterns automatically - Adapts to new patterns

#### 1.3.2 Why Machine Learning Now?

Three factors came together: 1. **Data** - We generate billions of data points daily 2. **Computing** - GPUs can process massive datasets 3. **Algorithms** - Better learning methods discovered

---

### 1.4 Part 2: The Three Learning Styles

#### 1.4.1 1. Supervised Learning (Learning with a Teacher)

**Example:** Spam detection - **Input:** Email text - **Label:** "Spam" or "Not Spam" - **Goal:** Learn to classify new emails

**Real Applications:** - Email spam filters (Gmail) - Fraud detection (credit cards) - Medical diagnosis (X-ray analysis) - Price prediction (real estate)

**When to Use:** - You have labeled examples (input + correct answer) - You want to predict or classify new data - Pattern is consistent over time

### 1.4.2 2. Unsupervised Learning (Learning without a Teacher)

**Example:** Customer segmentation - **Input:** Customer purchase history - **Label:** None - **Goal:** Discover natural customer groups

**Real Applications:** - Customer segmentation (marketing) - Anomaly detection (fraud, errors) - Recommendation systems (Netflix genres) - Document organization (Google News topics)

**When to Use:** - No labels available - Want to discover hidden patterns - Explore data structure

### 1.4.3 3. Reinforcement Learning (Learning by Trial and Error)

**Example:** Game playing - **Input:** Game state - **Action:** Make a move - **Reward:** Win/loss/draw - **Goal:** Learn winning strategy

**Real Applications:** - Game AI (AlphaGo, chess) - Robotics (walking, grasping) - Self-driving cars (navigation) - Resource optimization (data centers)

**When to Use:** - Sequential decisions matter - Learn from interaction - Delayed rewards

---

## 1.5 Part 3: How Does Learning Work?

### 1.5.1 The Learning Process

1. **Collect Data** - Gather examples (emails, images, sensor readings)
2. **Choose Model** - Pick learning algorithm (linear, tree, neural network)
3. **Train** - Show model examples, adjust internal parameters
4. **Evaluate** - Test on NEW data (not training data!)
5. **Deploy** - Use in real world

### 1.5.2 Key Insight: Generalization

**Goal:** Perform well on NEW data, not just training data

**Bad:** Memorizing training data (overfitting) **Good:** Learning underlying pattern (generalizing)

**Analogy:** Student preparing for exam - Memorizing exact practice problems = overfitting - Understanding concepts = generalizing

---

## 1.6 Part 4: Success Stories

### 1.6.1 Email Spam Detection

- **Before ML:** Rule-based filters caught 60% of spam
- **After ML:** Gmail catches 99.9% of spam
- **Why:** Learns new spam patterns automatically

### 1.6.2 Image Recognition

- **Before ML:** Hand-coded rules, poor accuracy
- **After ML:** 95%+ accuracy on ImageNet
- **Breakthrough:** Deep learning (2012)

### 1.6.3 Language Translation

- **Before ML:** Rule-based translation, awkward output
  - **After ML:** Google Translate uses neural networks
  - **Result:** Human-level quality for many language pairs
- 

## 1.7 Part 5: When NOT to Use Machine Learning

### 1.7.1 ML is NOT the answer when:

1. **Simple rules work** - Don't use ML to add two numbers
2. **No data available** - Need thousands+ examples minimum
3. **Explainability critical** - Medical/legal may require transparent logic
4. **Data changes rapidly** - Model becomes outdated quickly
5. **Cost exceeds benefit** - Training can be expensive

### 1.7.2 Traditional Programming is Better For:

- Calculations (tax computation)
  - Exact logic (password validation)
  - Known algorithms (sorting, searching)
  - Zero-tolerance errors (banking transactions)
- 

## 1.8 Part 6: Common Pitfalls (What Can Go Wrong)

### 1.8.1 1. Not Enough Data

**Problem:** Model can't learn pattern from 10 examples **Solution:** Collect more data (thousands minimum)

### 1.8.2 2. Biased Data

**Problem:** Training on non-representative data **Example:** Facial recognition trained only on light-skinned faces **Solution:** Diverse, balanced datasets

### 1.8.3 3. Overfitting

**Problem:** Model memorizes training data, fails on new data **Symptom:** 100% training accuracy, 50% test accuracy **Solution:** Regularization, more data, simpler model

### 1.8.4 4. Wrong Metric

**Problem:** Optimizing accuracy when precision matters **Example:** Cancer detection needs high recall (catch all cancers) **Solution:** Choose metric matching business goal

### 1.8.5 5. Data Leakage

**Problem:** Test data accidentally in training set **Result:** Falsely optimistic performance **Solution:** Strict train/test separation

---

## 1.9 Part 7: Practical Checklist

### 1.9.1 Before Starting ML Project:

- Do I have enough labeled data? (1000+ examples minimum)
- Is there a pattern to learn? (not random noise)
- Can I measure success clearly? (accuracy, profit, etc.)
- Is data representative of real-world use?
- Do I have computational resources? (GPU for deep learning)
- Is model interpretability required?
- What happens if prediction is wrong? (risk assessment)

### 1.9.2 Good First ML Projects:

1. **Classification:** Email spam, sentiment analysis
2. **Regression:** Price prediction, demand forecasting
3. **Clustering:** Customer segmentation, document grouping

### 1.9.3 Avoid As First Project:

- Real-time systems (latency critical)
  - Safety-critical applications (medical, automotive)
  - Highly imbalanced data (fraud: 0.1% positive rate)
  - Complex sequential decisions (reinforcement learning)
- 

## 1.10 Key Takeaways

1. **Learning from Data:** ML discovers patterns automatically vs hand-coded rules
  2. **Three Styles:** Supervised (with labels), Unsupervised (find patterns), Reinforcement (trial and error)
  3. **Generalization:** Goal is new data performance, not memorization
  4. **When to Use:** Abundant data, clear patterns, measurable outcomes
  5. **When NOT to Use:** Simple rules work, no data, explainability critical
  6. **Common Pitfalls:** Insufficient data, overfitting, bias, wrong metric
- 

## 1.11 Next Steps

- **Week 00b:** Supervised Learning algorithms (regression, trees, ensembles)
  - **Hands-On:** Try scikit-learn tutorials (spam detection, iris classification)
  - **Reading:** “Machine Learning Yearning” by Andrew Ng (free PDF)
- 

## 1.12 Glossary (Plain English)

- **Algorithm:** Step-by-step learning procedure
- **Feature:** Input variable (age, income, word count)
- **Label:** Correct answer for training example
- **Model:** Learned pattern from data
- **Training:** Process of learning from data
- **Testing:** Evaluating on new data

- **Overfitting:** Memorizing instead of learning
- **Generalization:** Working well on new data
- **Supervised:** Learning with labeled examples
- **Unsupervised:** Finding patterns without labels