

Neural Networks - Advanced Handout

Machine Learning for Smarter Innovation

1 Neural Networks - Advanced Handout

Target Audience: Data scientists and ML engineers **Duration:** 90 minutes reading **Level:** Advanced (mathematical foundations, optimization theory)

1.1 Mathematical Foundations

1.1.1 Forward Propagation

For layer l with weights $W^{(l)}$, bias $b^{(l)}$, and activation f :

$$\begin{aligned} z^{(l)} &= W^{(l)}a^{(l-1)} + b^{(l)} \\ a^{(l)} &= f(z^{(l)}) \end{aligned}$$

Where: - $a^{(0)} = x$ (input) - $a^{(L)} = \hat{y}$ (output)

1.1.2 Backpropagation

Loss gradient with respect to output:

$$\delta^{(L)} = \nabla_{a^{(L)}} \mathcal{L} \odot f'(z^{(L)})$$

Gradient propagation (for $l = L-1, \dots, 1$):

$$\delta^{(l)} = (W^{(l+1)})^T \delta^{(l+1)} \odot f'(z^{(l)})$$

Weight gradients:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial W^{(l)}} &= \delta^{(l)} (a^{(l-1)})^T \\ \frac{\partial \mathcal{L}}{\partial b^{(l)}} &= \delta^{(l)} \end{aligned}$$

1.2 Activation Functions

1.2.1 ReLU and Variants

ReLU: $f(x) = \max(0, x)$, $f'(x) = \mathbf{1}_{x>0}$

Leaky ReLU: $f(x) = \max(\alpha x, x)$ where $\alpha \approx 0.01$

GELU (Gaussian Error Linear Unit):

$$f(x) = x \cdot \Phi(x) \approx 0.5x(1 + \tanh[\sqrt{2/\pi}(x + 0.044715x^3)])$$

Swish: $f(x) = x \cdot \sigma(\beta x)$

1.2.2 Softmax for Classification

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Numerical stability:

$$\text{softmax}(z_i) = \frac{e^{z_i - \max(z)}}{\sum_{j=1}^K e^{z_j - \max(z)}}$$

1.3 Loss Functions

1.3.1 Cross-Entropy Loss

Binary:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Multi-class:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic})$$

1.3.2 Focal Loss (for imbalanced data)

$$\mathcal{L}_{FL} = -\alpha_t (1 - p_t)^\gamma \log(p_t)$$

Where γ is focusing parameter (typically 2).

1.4 Optimization Algorithms

1.4.1 Stochastic Gradient Descent (SGD)

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t; x_i, y_i)$$

1.4.2 SGD with Momentum

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} \mathcal{L}$$

$$\theta_{t+1} = \theta_t - v_t$$

1.4.3 Adam (Adaptive Moment Estimation)

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \\ \theta_{t+1} &= \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \end{aligned}$$

Default: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$

1.4.4 AdamW (Weight Decay Decoupled)

$$\theta_{t+1} = \theta_t - \eta \left(\frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} + \lambda \theta_t \right)$$

1.5 Initialization Strategies

1.5.1 Xavier/Glorot Initialization

For layer with n_{in} inputs and n_{out} outputs:

$$W \sim \mathcal{U} \left(-\sqrt{\frac{6}{n_{in} + n_{out}}}, \sqrt{\frac{6}{n_{in} + n_{out}}} \right)$$

For tanh activation. Maintains variance through layers.

1.5.2 He/Kaiming Initialization

$$W \sim \mathcal{N} \left(0, \sqrt{\frac{2}{n_{in}}} \right)$$

For ReLU activation. Accounts for ReLU zeroing half the activations.

```
# PyTorch implementation
nn.init.kaiming_normal_(layer.weight, mode='fan_in', nonlinearity='relu')
nn.init.xavier_uniform_(layer.weight)
```

1.6 Batch Normalization

1.6.1 Forward Pass

$$\begin{aligned} \hat{x}_i &= \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \\ y_i &= \gamma \hat{x}_i + \beta \end{aligned}$$

Where μ_B , σ_B^2 are batch statistics; γ , β are learned.

1.6.2 Inference

Use running averages instead of batch statistics:

$$\mu_{running} = (1 - \alpha)\mu_{running} + \alpha\mu_B$$

1.6.3 Layer Normalization (for sequences)

Normalize across features instead of batch:

$$\hat{x}_i = \frac{x_i - \mu_L}{\sqrt{\sigma_L^2 + \epsilon}}$$

1.7 Regularization Theory

1.7.1 Dropout

During training, randomly zero activations with probability p :

$$\tilde{a} = \frac{1}{1-p} \cdot a \cdot m, \quad m_i \sim \text{Bernoulli}(1-p)$$

Interpretation: Ensemble of 2^n sub-networks.

1.7.2 L2 Regularization (Weight Decay)

$$\mathcal{L}_{total} = \mathcal{L}_{data} + \frac{\lambda}{2} \|W\|_2^2$$

Gradient becomes:

$$\nabla_W \mathcal{L}_{total} = \nabla_W \mathcal{L}_{data} + \lambda W$$

1.7.3 Label Smoothing

Instead of one-hot targets, use:

$$y_{smooth} = (1 - \alpha)y_{one-hot} + \frac{\alpha}{K}$$

Typically $\alpha = 0.1$.

1.8 Attention Mechanism

1.8.1 Scaled Dot-Product Attention

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Where: - $Q \in \mathbb{R}^{n \times d_k}$ (queries) - $K \in \mathbb{R}^{m \times d_k}$ (keys) - $V \in \mathbb{R}^{m \times d_v}$ (values)

1.8.2 Multi-Head Attention

$$\begin{aligned} \text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \\ \text{head}_i &= \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \end{aligned}$$

1.8.3 Self-Attention

When $Q = K = V = X$ (same input):

$$\text{SelfAttention}(X) = \text{softmax} \left(\frac{XX^T}{\sqrt{d}} \right) X$$

1.9 Transformer Architecture

1.9.1 Encoder Block

```
Input -> LayerNorm -> MultiHeadAttention -> Residual -> LayerNorm -> FFN ->
Residual -> Output
```

1.9.2 Position Encoding

$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

1.9.3 Implementation

```
class TransformerBlock(nn.Module):
    def __init__(self, d_model, n_heads, d_ff, dropout=0.1):
        super().__init__()
        self.attention = nn.MultiheadAttention(d_model, n_heads, dropout=dropout)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.ff = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.GELU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model),
            nn.Dropout(dropout)
        )

    def forward(self, x, mask=None):
        # Self-attention with residual
        attn_out, _ = self.attention(x, x, x, attn_mask=mask)
        x = self.norm1(x + attn_out)

        # Feed-forward with residual
        ff_out = self.ff(x)
        x = self.norm2(x + ff_out)

    return x
```

1.10 Gradient Issues

1.10.1 Vanishing Gradients

Cause: Saturating activations (sigmoid, tanh) or deep networks.

Solutions: - ReLU activations - Batch/Layer normalization - Residual connections - LSTM/GRU for sequences

1.10.2 Exploding Gradients

Cause: Large weight magnitudes, especially in RNNs.

Solutions: - Gradient clipping: $g \leftarrow \min(1, \frac{\theta}{\|g\|}) \cdot g$ - Weight initialization - Layer normalization

```
# Gradient clipping in PyTorch
torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
```

1.11 Learning Rate Scheduling

1.11.1 Cosine Annealing

$$\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min})(1 + \cos(\frac{t}{T}\pi))$$

1.11.2 Warmup + Decay

$$\eta_t = \begin{cases} \eta_{max} \cdot \frac{t}{T_{warmup}} & t < T_{warmup} \\ \eta_{max} \cdot \text{decay}(t - T_{warmup}) & t \geq T_{warmup} \end{cases}$$

```
# PyTorch scheduler
scheduler = optim.lr_scheduler.CosineAnnealingWarmRestarts(
    optimizer, T_0=10, T_mult=2
)

# OneCycleLR (recommended for training)
scheduler = optim.lr_scheduler.OneCycleLR(
    optimizer, max_lr=0.01, epochs=epochs, steps_per_epoch=len(train_loader)
)
```

1.12 Mixed Precision Training

1.12.1 FP16 Training

```
from torch.cuda.amp import autocast, GradScaler

scaler = GradScaler()

for batch in train_loader:
    optimizer.zero_grad()

    with autocast():
        outputs = model(inputs)
        loss = criterion(outputs, targets)

    scaler.scale(loss).backward()
    scaler.step(optimizer)
```

```
scaler.update()
```

Benefits: 2x memory reduction, faster training on modern GPUs.

1.13 Distributed Training

1.13.1 Data Parallel

```
model = nn.DataParallel(model) # Simple, single-machine multi-GPU
```

1.13.2 Distributed Data Parallel

```
import torch.distributed as dist
from torch.nn.parallel import DistributedDataParallel as DDP

dist.init_process_group("nccl")
model = DDP(model, device_ids=[local_rank])
```

1.14 Model Compression

1.14.1 Quantization

```
# Post-training quantization
quantized_model = torch.quantization.quantize_dynamic(
    model, {nn.Linear}, dtype=torch.qint8
)
```

1.14.2 Knowledge Distillation

$$\mathcal{L}_{KD} = \alpha \mathcal{L}_{CE}(y, \hat{y}_{student}) + (1 - \alpha) T^2 \mathcal{L}_{KL}(\sigma(z_T/T), \sigma(z_S/T))$$

1.14.3 Pruning

```
import torch.nn.utils.prune as prune

# Prune 30% of weights with smallest magnitude
prune.l1_unstructured(module, name='weight', amount=0.3)
```

1.15 References

1. Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.
2. Vaswani, A., et al. (2017). “Attention Is All You Need”

-
- 3. He, K., et al. (2016). “Deep Residual Learning for Image Recognition”
 - 4. Kingma, D. P., & Ba, J. (2015). “Adam: A Method for Stochastic Optimization”
-

Deep learning is an empirical science. Theory provides guidance, but experimentation determines success.